



TikTok's DSA Transparency Report

July - December 2025



Executive Summary

At TikTok, our mission is to inspire creativity and bring joy. The safety and well-being of our community is our priority. TikTok has a strong track record in voluntary transparency reporting; we have been publishing [transparency reports](#) since 2019. We also report on our efforts to combat disinformation on our platform under the [Code of Conduct on Disinformation](#) on a bi-annual basis. Building on our transparency efforts and in line with our obligations under the Digital Services Act (DSA), we are pleased to publish our sixth DSA transparency report, in the template provided by the Implementing Regulation, for the reporting period of 1 July 2025 to 31 December 2025. The required template can be found in the attached files. The following document provides additional information, outlining methodological choices, contextual information and a summary.

We have a number of measures designed to keep users safe across priority areas, including from illegal and other harmful content. We are pleased to report on the numbers underlying these measures including the additional reporting option we have implemented to allow people to report content in the European Union they believe is illegal. Key insights from this report include:

- **Enforcing our rules by removing violative content:** During this reporting period, we have removed around **112 million pieces of content** that violated our Policies. These metrics record the number of content removals across Video, Live streams, Ads, Product Listings and, for the first time, the number of Comments that were removed.
- **Automation is increasingly effective at improving speed and consistency of moderation at scale.** Automated systems actioned 93.8% of all violating content without human review. These decisions maintained a high degree of precision, with 97.6% of automated enforcement decisions being confirmed as correct.

Providing transparency to our community about how we keep them safe is an ongoing commitment that has no finish line. We are pleased to have built on the efforts of our previous DSA transparency reports, for example, **by adding enforcement volumes on comments.**



Report index

- 1. Executive summary**
- 2. Categories of illegal and incompatible content**
- 3. Member State orders to act against illegal content and provide information**
- 4. Notices submitted in accordance with Article 16 DSA**
- 5. Own-initiative content moderation on the basis of illegality**
- 6. Own-initiative content moderation on the basis of Terms of Service**
- 7. Internal complaints mechanism, out-of-court dispute settlement bodies and suspensions imposed on repeated offenders**
- 8. Use of automated means for content moderation and indicators of accuracy**
- 9. Human resources dedicated to content moderation**
- 10. Average monthly active recipients**



2. Categories of illegal and incompatible content

For this reporting period, and to reflect the requirements of the transparency reporting template, TikTok conducted a comprehensive review and mapping of its moderation labels to align them with the categories set out in the Implementing Regulation.

The use of Policies in this report refers collectively to our applicable terms and policies, including our [Terms of Service](#), [Community Guidelines](#), [Advertising Policies](#) and [TikTok Shop Policies](#) (together, our Policies).

This mapping follows the same categorisation framework TikTok used when submitting Statements of Reasons to the DSA Transparency Database, ensuring consistency across its regulatory reporting.

As a result of this alignment, the category-level figures in this transparency report may differ from those published in previous reporting periods, which were based on TikTok's internal policy taxonomy. More information on the mapping is available in the **"2_categories_names"** tab of the transparency reporting template.

In addition, TikTok has updated its method for counting enforcement actions to comply with the European Commission's guidance on avoiding double counting in the template.

In previous reports, a single piece of content that was removed for multiple policy violations was counted once at the aggregate level, but also counted separately under each violated policy. As a result, the policy-level totals were higher than the overall enforcement figure.

Under the new methodology, TikTok spreads policy violations proportionally across the total number of content actioned, without counting anything twice. This means the numbers in each category add up exactly to the overall total, while still reflecting how our rules are enforced in practice.

3. Member State orders to act against illegal content and provide information

When we receive requests from government authorities in the European Union to remove content, we evaluate in line with our Policies and applicable law. When we receive requests for user information disclosure, our response is developed with the [freedom of expression and privacy](#) of our users in mind. All requests are reviewed individually in line with our [Law Enforcement Guidelines](#) and we only disclose user data where a request is legitimate and based on a valid legal process.

During the period 1 July 2025 to 31 December 2025, TikTok received orders from government authorities in the European Union to both remove content and disclose user information. The tab **"3_member_states_orders"** sets out the number of orders received, the number of items identified in each order, and median turnaround times, broken down by Member State market and category of illegal content.

Member State orders to act against illegal content

TikTok has a dedicated channel through which government authorities in the European Union may submit orders to remove content. We confirm receipt by sending an acknowledgement. The number of orders received, items per order, median acknowledgement times, and median times to give effect to orders are set out in the tab **"3_member_states_orders"**, broken down by Member State market and category of illegal content. The median time to give effect to an order reflects the period between receipt of a valid order and TikTok investigating and actioning it — either by removing the content or providing a substantive response to the issuing authority.

Member State orders to provide information

Government authorities in the European Union may also submit orders for disclosure of information through TikTok's dedicated channel. We acknowledge receipt of an order from a government authority submitted through our



dedicated channel immediately, by sending an automatic acknowledgement. The number of orders received, items per order, median acknowledgement times, and median times to give effect to orders are set out in tab **"3_member_states_orders"**, broken down by Member State market and category of illegal content. Median turnaround times reflect that certain orders involve multiple pieces of content, accounts, or identifiers requiring additional review, as well as cases where clarification is requested from the issuing authority but no response is received. Those cases are closed after 28 days.

Methodological Note

For this reporting period, TikTok has expanded the scope of metrics reported in relation to government orders, in line with the new transparency reporting template. TikTok now reports (i) the number of individual content items identified in each request ticket, and (ii) median response times disaggregated by both Member State market and illegal content category. In prior reporting periods, median response times were reported only at the aggregate level.

In addition, TikTok has revised its approach to the categorisation of Member State government orders. As government authorities are not required to select a specific content type when submitting orders, all tickets received during this period have been classified as "Type of illegal content not specified by the public authority." This ensures the reported data faithfully represents the information as submitted and avoids interpretive categorisation by TikTok's internal teams.

4. Notices submitted in accordance with Article 16 DSA

Our Policies apply to all accounts and content on the platform. They often align with, and sometimes go beyond, local law requirements. While we primarily enforce our Policies at our own initiative through automated and human moderation, users can also use the reporting functions to alert TikTok to content they believe violates our Policies. The number of user reports made in the European Union to TikTok during the period 1 July 2025 to 31 December 2025 is set out in tab **"4_notices"**. Under the DSA, "trusted flaggers" can also submit illegal content reports.

As part of our requirements under the DSA, we introduced an [additional reporting channel](#) for our community in the European Union to 'Report Illegal Content', which enables users to alert us to content they believe breaches the law. When users report suspected illegal content, they will be asked to select a category of illegal content they are reporting under. Reporters are also asked to provide additional information, such as: the jurisdiction in question; if possible, the specific law in question; and a clear explanation as to why they think the content violates the law. If the report is incomplete (for example, it does not provide enough information for us to assess if the content is illegal) or materially unsubstantiated, the report may be rejected. The reporter will be notified of this decision and provided with an opportunity to re-submit their report with more information. This helps us properly and effectively consider and respond to each report.

Illegal content reports are processed using a combination of automation and human moderation. We first assess the reported content against our Policies. If a violation is identified, the content may be removed globally. If the content does not violate our Policies, our moderation team focusing on illegal content conducts a further human review to determine whether it is unlawful in the relevant jurisdiction. This assessment includes balancing competing legal rights, such as freedom of expression. We notify the reporting party of our decision, including where we determine the content is not illegal. Users who disagree with a decision may [appeal](#) through our appeals process.

Methodological Note

In accordance with the requirements of the new reporting template, TikTok has enhanced the granularity of the metrics reported in relation to notices received from users and trusted flaggers. Specifically, median response times are now reported at a more granular level than in prior reporting periods. Readers should note that, due to technical constraints associated with aggregating response time data across multiple product features and underlying datasets, median response times for notices are reported only at the subcategory level rather than the broader category level.



5. Own-initiative content moderation on the basis of illegality

We assess the legality of content where it is reported to us as suspected illegal content, including through user reports, Trusted Flagging notices, or government orders. Outside of these channels, our proactive detection efforts focus on identifying and enforcing violations of our Policies.

6. Own-initiative content moderation on the basis of Terms of Service

TikTok strives to foster an open and inclusive environment where people can create, find community, and be entertained. To maintain that environment, we take action upon content and accounts that violate our Policies. We are committed to being transparent with our community about the moderation actions we take.

The number and type of restrictions we impose as part of our content moderation activities are available in tab **"6_own_initiative_TC"**.

Our Policies are the starting point when it comes to how we form and operate our content moderation strategies and practices and they contain provisions which prohibit various forms of illegal and other harmful content. We use a combination of automation and human moderation to identify, review, and action content that violates our Policies.

Methodological Note

For the purposes of this reporting period, TikTok has introduced several methodological changes to the reporting of its own-initiative content moderation activities undertaken in respect of violations of its Policies.

First, enforcement metrics that were previously reported separately for Trust & Safety, Ads, and TikTok Shop are now aggregated across all business organisations and mapped to the template's illegal content categories. This alignment ensures consistency with the categorisation framework described in Section 2 of this report and facilitates comparability across the various reporting sections.

Second, metrics relating to content demotion have been refined. In prior transparency reports, the "content demoted" category encompassed both age-restriction measures and feed-eligibility restrictions. Going forward, age restrictions are reported as a distinct metric, separate from content demotion, to provide greater transparency regarding the nature of the enforcement action taken.

Third, the double counting avoidance methodology described in Section 2 applies equally to the own-initiative enforcement figures reported in this section. The proportional distribution approach ensures that the sum of subcategory-level figures aligns exactly with the total enforcement count, while maintaining consistency with TikTok's actual enforcement patterns.

Finally, TikTok has expanded the scope of content types covered in this reporting section. In addition to own-initiative enforcement data on videos, LIVE streams, advertisements, and product listings, this report now also includes enforcement actions taken against comments, reflecting a broader and more comprehensive view of TikTok's content moderation activities across the platform.

7. Internal complaints mechanism, out-of-court dispute settlement bodies and suspensions imposed on repeated offenders

Complaints

Anyone can report content on TikTok they believe violates our Policies or applicable laws and can appeal if they disagree with the outcome of our decision. Users are also provided with a notification where we determine that they have violated our Policies or applicable laws, and they are provided with an opportunity to appeal against any action we have taken. We provide information to users about how to appeal a decision in relation to a report they have made, or how to appeal a decision affecting their content or account [here](#). We report on the number of appeals, and the action we take in response to those appeals, in tab **"7_appeals_and_recidivism"**.



Disputes submitted to out-of-court dispute settlement bodies

The DSA provides users of the platform with the right to access a third party out-of-court dispute settlement process to resolve any disputes that they may have regarding moderation actions (including in relation to any appeals), essentially allowing users to submit appeals of content moderation decisions to an independent third party for assessment. We report on the number of appeals submitted to these bodies, the outcomes of the dispute settlement, the median time needed for completing the dispute settlement procedures, and the share of appeals where we implemented the decisions of the body, in tab "**7_appeals_and_recidivism**".

8. Use of automated means for content moderation and indicators of accuracy

Key principles

We operate our content moderation processes using a combination of automation and human moderation in accordance with the following four pillars, which provide that we will:

1. Remove content from the platform that violates our Policies;
2. Age-restrict mature content (that does not violate our Community Guidelines but which contains mature themes) so it is only viewed by adults (18 years and older);
3. Maintain For You feed eligibility standards to help ensure any content that may be promoted by the recommendation system is appropriate for a broad audience; and
4. Empower our community with information, tools, and resources.

Automated Review

We place considerable emphasis on proactive detection to remove violative content and reduce exposure to potentially distressing content for our human moderators. Before content is posted to our platform, it's reviewed by automated moderation technologies which identify content or behavior that may violate our policies or For You feed eligibility standards, or that may require age-restriction or other actions. While undergoing this review, the content is visible only to the uploader.

If our automated moderation technology identifies content that is a potential violation, it will either take action against the content or flag it for further review by our human moderation teams. In line with our safeguards to help ensure accurate decisions are made, automated removal is applied when violations are the most clear-cut.

Some of the methods and technologies that support these efforts include:

- **Vision-based:** Computer vision models can identify objects that violate our Policies, such as weapons or hate symbols.
- **Audio-based:** Audio clips are reviewed for violations of our Policies, supported by a dedicated audio bank and "classifiers" that help us detect audios that are similar or modified to previous violations.
- **Text-based:** Detection models review written content like comments or hashtags, using foundational keyword lists to find variations of violative text. Artificial Intelligence (AI) that can interpret the context surrounding content—helps us identify violations that are context-dependent, such as words that can be used in a hateful way but may not violate our policies by themselves. We also work with various external experts, like our [fact-checking partners](#), to inform our keyword lists.
- **Similarity-based:** "Similarity detection systems" enable us to not only catch identical or highly similar versions of violative content, but other types of content that share key contextual similarities and may require additional review.
- **Activity-based:** Technologies that look at how accounts are being operated help us disrupt deceptive activities like bot accounts, spam, or attempts to artificially inflate engagement through fake likes or follow attempts.
- **LLMs:** We use multimodal LLMs to help moderate content faster and more consistently at scale, from taking automated action on activity like fake engagement, to empowering teams with better moderation tools and risk insights.

- We work with external groups, for example [Tech Against Terrorism](#) in the context of [violent extremist](#) content, who help us to more quickly detect and remove violative content that has already been identified off the platform.

Ongoing improvements

Throughout 2025, we have been making ongoing improvements to our safety technologies and enhancements to our content moderation processes. These investments continue to bear out. For example, globally in Q3 2025, we shared our:

- Highest-ever rate of proactive removal rate (content removed before it's reported to us): 99.3%
- Highest-ever volume of content removed in under 24 hours: 94.8%
- Highest-ever rate of violative content removed by automated technologies: 91%

If users or advertisers believe we have made a mistake, they can appeal the removal of their content.

These advances are helping us remove violative content more quickly, reducing the likelihood of our community seeing it. As more of these tasks are handled by technology, our safety teams can focus more of their time on work that most benefits from human expertise—such as handling appeals, consulting external experts, and responding to fast-moving events.

In addition, continued improvements in AI and moderation technologies support the well-being of our safety teams by reducing their exposure to potentially distressing content and equipping them with better tools to carry out this critical work effectively. For example, technologies like AI help make it easier to moderate nuanced areas like misinformation by extracting specific misinformation "claims" from videos for moderators to assess directly or route to our fact-checking partners.

Methodological Note

For this reporting period, TikTok has adopted a revised methodology for assessing the performance of its automated content moderation technologies, in line with the template's requirement in tab "**8_automated_means**" to report proxy metrics for precision, recall, and accuracy. The three indicators are defined as follows:

- **Precision** measures the correctness of the automated system when it takes enforcement action. It is calculated as the proportion of automated enforcement actions that were upheld (i.e., not successfully appealed) out of all automated enforcement actions during the reporting period.
- **Recall** measures the completeness of automated detection. It is calculated as upheld automated enforcement actions divided by total upheld enforcement actions across both automated and human detection.
- **Accuracy** provides an overall measure of correctness across automated enforcement decisions. It is calculated as upheld automated enforcement actions plus restored human enforcement actions, divided by total enforcement outcomes.

This framework relies on two principal ground-truth inputs:

1. **User and advertisers appeals**, which serve as an indicator of whether an automated enforcement action was correct or erroneous; and
2. **Human review** outcomes, which capture instances of content that the automated system did not detect but which were subsequently identified and actioned by human moderators.

TikTok acknowledges that these inputs function as **proxy indicators** and that the resulting metrics are subject to the inherent limitations of the underlying data sources. TikTok continues to refine its measurement capabilities and will update its methodology as improved data and analytical frameworks become available.



9. Human resources dedicated to content moderation

Our Trust & Safety teams lead our approach to content moderation. They are responsible for the development of our Community Guidelines and related moderation policies, and for the moderation of content on TikTok. For the European Union, TikTok's Trust & Safety work is led from Dublin, Ireland, where a number of key global Trust & Safety personnel are based. For ads and branded content, our Monetisation Integrity and Business Integrity teams are responsible for TikTok's Advertising Policies and related moderation policies, and play a key role in their moderation. Our TikTok Shop Governance & Experience teams develop TikTok Shop policies, and are involved in the moderation of TikTok Shop listings and other TikTok Shop-related content.

Training

To ensure a consistent understanding and application of our Policies, all content moderator personnel receive training across our relevant Policies. All content moderators undergo training on TikTok's content moderation systems and moderator well-being. Personnel involved in reviewing reported illegal content receive additional focused training on assessing the legality of reported illegal content.

Content moderation training materials are kept under review to ensure that they are accurate and current. Such materials include clearly defined learning objectives to ensure our content moderators understand the core policy issues and their underlying policy rationale, key terms and policy exceptions (where applicable).

Members of our safety teams attend regular internal sessions dedicated to knowledge sharing and discussion about relevant issues and trends, which include input from external experts. For example, as part of our Election Speaker Series, which helps inform our approach to elections, we invite suitably qualified external local and regional experts to share their insights and market expertise with our internal teams. During the reporting period, we ran 4 Election Speaker Series sessions, 3 in EU Member States, Czechia, Ireland, and Netherlands, and 1 in Moldova.

Our teams also participate in various external events to share expertise and support their continued professional learning. These engagements contribute to the team's awareness of the risks which may arise on the platform, which in turn informs our approach to content moderation.

Support

Human safety professionals continue to play a crucial role in our content moderation approach. We have thousands of safety professionals globally who help build our technologies, develop and enforce our policies, design new safety features, and work with experts and industry peers. We strive to promote a caring working environment for all TikTok employees, and especially for trust and safety professionals. We use an evidence-based approach to develop programs and resources that support their psychological well-being.

Our primary focus is on preventative care measures to minimise the risk of psychological injury through well-timed support, training and tools, from recruitment through to onboarding and ongoing employment, that help foster resilience while minimising the risk of psychological injury. These may include tools and features to allow employees to control exposure to graphic content when reviewing or moderating content, including grayscaling, muting and blurring; training for managers to help them identify when a team member may need additional well-being support; and clinical and therapeutic support. We also continue to lean into moderation technology as an effective way to reduce human moderators' exposure to harmful content.

We provide our trust and safety employees with membership to the Trust and Safety Professional Association (TSPA). This membership allows them to access resources for career development, participate in workshops and events, and connect with a network of peers across the industry.

Qualifications & linguistic expertise

Some of the issues which arise on the platform are highly localised in terms of language and region, which requires deep knowledge and awareness of relevant cultural nuances, terms and context.

To address this, and ensure its content moderators are appropriately qualified to make decisions, we have regional policy teams in each region, which includes coverage for all European Union Member States, for example with either designated policy country managers for larger countries or policy managers covering a number of smaller countries.



Based primarily in Europe, Regional Policy teams bring regional insights, cultural context, and local expertise to ensure that global moderation policies are applied appropriately across diverse countries and communities in Europe. Acting as expert policy ambassadors, they work to create a safe and positive experience for users by ensuring that our Policies are upheld in ways that reflect local realities. They play a key role in risk mitigation by detecting and addressing regional trends, engaging with external experts such as NGOs, civil society organisations, and government authorities, and collaborating closely with cross-functional teams. Their work includes developing policy interventions and enforcement strategies that strengthen our ability to reduce harm, while maintaining a safe and welcoming environment for our community.

The localised policy outputs from the EMEA regional policy team enable our content moderation teams to take a regionally informed approach to content moderation (e.g. rapidly evolving alternative vocabulary or terminology in relation to an unfolding election issue, which may vary/evolve over time and as between countries and languages).

We have also established a number of specialised moderation teams to assist our moderators to review content relating to complex issues. For example, assessing harmful misinformation requires additional context and assessment by our misinformation moderators who have enhanced training, expertise and tools to identify such content, including our global repository of previously fact-checked claims from the IFCN-accredited fact-checking partners and direct access to our fact-checking partners where appropriate.

We moderate content in more than 70 languages globally and we are transparent in our [regular Community Guidelines Enforcement Reports](#) about the primary languages our moderators work in globally. We have language capabilities covering at least one official language for each of the 27 European Union Member States, as well as a number of other languages that are commonly spoken in the region (for example, Arabic and Turkish). This language capability complements our awareness-raising materials, like the Community Guidelines, that are also available in multiple languages. We also have moderation personnel that are not assigned to a particular language, who assist with reviewing content such as photos and profiles.

Human moderation

In order to support fair and consistent review of potentially violative content, our moderators work alongside our automated moderation systems and take into account additional context and nuance which may not always be picked up by technology.

Human moderation also helps improve our automated moderation systems by providing feedback for the underlying machine learning models to strengthen our ongoing detection capabilities. This continuous improvement helps to reduce the volume of potentially distressing videos that moderators view and enables them to focus more on content that requires a greater understanding of context and nuance (such as misinformation, hate speech and harassment). The responsibilities of our content moderators include:

- **Reviewing content flagged by technology:** When our automated moderation systems identify potentially problematic content but cannot make an automated decision to remove it, they send the content to our moderation teams for further review. To support this work, we have developed technology that can identify potentially violative items – for example, emblems associated with extremist groups – in video frames, so that content moderators can carefully review the video and the context in which it appears. This technology improves the efficiency of moderators by helping them more adeptly identify violative images or objects, quickly recognise violations, and make decisions accordingly.
- **Reviewing reports from its community:** Community reporting helps us maintain a safe environment. We offer our community easy-to-use in-app and online reporting tools so they can flag any content or account they feel is in violation of our Policies or may be illegal. TikTok's trusted partners also play a role in helping detect and remove harmful content. Through our Community Partner Channel Program, participating civil society organizations with a broad range of expertise across online and content safety can report potentially violative or illegal content to us directly for review. Community reports are an important component of our content moderation process. However, the vast majority of removed content is identified proactively before



it is reported to us.

- **Reviewing popular content:** We manually review video content when it reaches certain levels of popularity in terms of the number of video views, reducing the risk of violative content being shown in the For You Feed or otherwise being widely disseminated.
- **Assessing appeals:** If someone disagrees with our decision to restrict or remove their content or an account, they can appeal the decision for reconsideration. These appeals may be sent to moderators to decide if the content should be reinstated on the platform or the account reinstated.

The tab **"9_human_resources"** sets out the number of people who are dedicated to content moderation in line with our Policies and applicable local laws, broken down per each of the official European Union languages. Our moderators often have linguistic expertise across multiple languages. Where our moderators have linguistic expertise in more than one European Union language, that expertise is reflected in the detailed language breakdown below. For example, the Czech, Slovakian and Slovenian languages are grouped under one category within our Trust & Safety team and are moderated by the same moderators. The moderators allocated for the Croatian language also cover the Serbian language. These numbers also include moderators covering a number of other languages that are commonly spoken in the region, such as Arabic, Catalan, Hindi, Pashto, Persian, Turkish, Ukrainian, Norwegian, Russian and Icelandic.

The numbers in tab **"9_human_resources"** do not reflect the broader teams who also play a key role in keeping our community safe (for example, those involved in the development of our content moderation policies). The total number of human moderators includes non-language specific moderators (meaning the moderators who review profiles or photos).

10. Average monthly active recipients

We report on the average number of 'monthly active recipients', broken down per each of the 27 European Union Member States in tab **"10_AMAR"** during the period from 1 July to 31 December 2025.

Tab **"10_AMAR"** sets out the average number of 'monthly active recipients' in the European Union broken down per each Member State during the period 1 July 2025 and 31 December 2025, rounded to the nearest hundred thousand.

We have produced this calculation for the purposes of complying with our obligations under Article 42(3) of the DSA and it should not be relied on for other purposes. We have applied the same methodology used when calculating our total monthly active recipients number for the European Union published in [February 2026](#). In light of our legal requirements to provide the number broken down per Member State and given that users may have accessed the platform from different Member States in the relevant period, the estimates below may mean, in certain limited circumstances, user access is counted more than once.

Where we have shared user metrics in other contexts, the methodology and scope may have differed. Our approach to producing this calculation may evolve or may require altering over time, for example, because of product changes or new technologies.