Academic Literature Search — "Moderation Dynamics in Transition" (Bachelor's Thesis, University of Paderborn)

TL;DR

- The most directly relevant prior work for this thesis clusters around four anchor studies — Trujillo, Fagni & Cresci (2025); Papaevangelou & Votta (2025); Drolsbach & Pröllochs (2024); and Kaushal et al. (2024) — which collectively establish the empirical templates, comparative TikTok-vs-X metrics, and methodological caveats the thesis must engage with.
- For RQ1 (modality / affordances), the Bucher & Helmond (2018) chapter should be paired with Bossetta (2018) on digital architectures and with Zeng & Kaye (2022) on TikTok's visibility moderation; for RQ2 (event-driven moderation), the strongest empirical templates are Shahi et al. (2025) on the 2024 EP election and Pierri et al. (2023) on Twitter during geopolitical events.
- The 2025 German Federal Snap Election context is best supported by emerging grey literature (DFRLab/AlgorithmWatch "Musk Effect"; CeMAS Bundestagswahl monitoring; DRI "Filtered for You"; ISD Germany; Global Witness; Solovev, Drolsbach, Demirel & Pröllochs 2025/2026) — published material that directly observed TikTok and X moderation/recommendation behaviour around 23 February 2025.

Key Findings

This review identifies 30+ studies across the seven thesis dimensions, with the densest coverage in DSA-TDB empirical work (Dimension 1) and platform observability (Dimension 5). Three notable gaps emerged: (a) almost no peer-reviewed work yet exists that uses DSA-TDB data specifically for the 2025 German federal election — this is a genuine research opening for the thesis; (b) agenda-setting (Dimension 7) is mostly addressed through pre-DSA frameworks that must be theoretically bridged to algorithmic governance contexts; and (c) the structural API v1→v2 schema break (July 2025) is acknowledged in policy literature (European Commission documentation and Husovec's DSA newsletters) but not yet in any peer-reviewed empirical study — a methodological contribution opportunity.

PART A — STUDIES GROUPED BY THE 7 DIMENSIONS

Dimension 1 — DSA Transparency Database: Empirical Studies

🟢 **Trujillo, A., Fagni, T., & Cresci, S. (2025).** "The DSA Transparency Database: Auditing Self-reported Moderation Actions by Social Media." *Proceedings of the ACM on Human-Computer Interaction* 9(2), Article CSCW187, 1–28. DOI: 10.1145/3711085. arXiv:2312.10269. **[2024–2026, EU context]** *Relevance:* Analyses all 353.12M records submitted by the eight largest social media platforms in the EU during the first 100 days of the database, with a platform-wise comparison of volume, grounds, restriction types, content types, timeliness and automation, plus cross-checks against the platforms' Article 15 reports. Directly provides the methodological template for the thesis's cross-platform comparison and identifies X as the platform with the most reporting inconsistencies. *Dimensions: 1, 2, 5, 6.* arXiv + 2

Drolsbach, C. P., & Pröllochs, N. (2024). "Content Moderation on Social Media in the EU: Insights From the DSA Transparency Database." *Companion Proceedings of the ACM Web Conference 2024 (WWW '24 Companion)*, 939–942. DOI: 10.1145/3589335.3651482. arXiv:2312.04431. **[2024–2026, EU context]** *Relevance:* First empirical analysis of 156M SoRs over the database's first two months. Reports the now-canonical headline that "TikTok performs more than 350 times more content moderation decisions per user than X/Twitter" and shows the predominance of automated moderation, with X as the main exception — both findings are foundational for the thesis's modality comparison. *Dimensions: 1, 2, 6.* arXiv + 2

Kaushal, R., Van de Kerkhof, J., Goanta, C., Spanakis, G., & Iamnitchi, A. (2024). "Automated Transparency: A Legal and Empirical Analysis of the Digital Services Act Transparency Database." In *FAccT '24: Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, 1121–1132. DOI: 10.1145/3630106.3658960. arXiv:2404.02894. **[2024–2026, EU context]** *Relevance:* Combines doctrinal legal analysis of the DSA's transparency obligations with empirical analysis of 131M SoRs from November 2023. Introduces the framing concept of "automated transparency" and finds that >99.8% of SoRs are based on ToS infringements rather than statutory illegal-content grounds — a critical caveat for interpreting any decision-grounds distribution in the thesis. *Dimensions: 1, 5.* arXiv + 3

Dergacheva, D., Kuznetsova, V., Scharlach, R., & Katzenbach, C. (2023). "One Day in Content Moderation: Analyzing 24h of Social Media Platforms' Content Decisions through the DSA Transparency Database." Technical Report, Center for Media, Communication and Information Research (ZeMKI), University of Bremen / Lab Platform Governance, Media, and Technology (PGMT). URL: platform-governance.org/2023/one-day-in-content-moderation-by-social-media-platforms-in-the-eu/. **[EU context]** *Relevance:* Earliest empirical snapshot of the DSA-TDB (~2M SoRs across 24h); foregrounds platform heterogeneity in content categories and the prevalence of automation. Provides a useful one-day baseline against which longer longitudinal panels (including the thesis) can be benchmarked. *Dimensions: 1, 5, 6.* arXiv

Shahi, G. K., Tessa, B., Trujillo, A., & Cresci, S. (2025). "A Year of the DSA Transparency Database: What it (Does Not) Reveal About Platform Moderation During the 2024 European Parliament Election." *Proceedings of the ICWSM Workshops 2025.* arXiv:2504.06976. **[2024–2026, EU context]** *Relevance:* Per the arXiv abstract: "we analyze 1.58 billion self-reported moderation actions taken by eight large social media platforms during an extended period of eight months surrounding the 2024 European Parliament elections" (note the paper body cites ten months — this internal inconsistency is worth flagging in the thesis). Although the thesis explicitly excludes the 2024 EP election as an analog, this study is methodologically central: it concludes that platforms did not exhibit significant moderation-strategy changes around the election. The thesis can engage critically with this finding — testing whether a *snap* national election shows a different pattern. *Dimensions: 1, 4, 5.* ResearchGate +2 ResearchGate +2

Tessa, B., Amram, D., Monreale, A., & Cresci, S. (2025). "Improving Regulatory Oversight in Online Content Moderation." *IAIL '25: Imagining the AI Landscape after the AI Act*, June 9, 2025, Pisa. arXiv:2506.04145. **[2024–2026, EU context]** *Relevance:* Proposes a "Transparency Report Cross-Checking Process" and a "Verification Process" to identify inconsistencies between DSA-TDB submissions and Article 15 reports. Directly informs the thesis's API v1/v2 harmonization strategy and the methodological treatment of self-reported data. *Dimensions: 1, 5.* Semantic Scholar arxiv

Groesch, S., Birrer, A., Just, N., & Saurwein, F. (2026). "Big data, small answers: How the DSA Transparency Database falls short of its regulatory objectives." *Telecommunications Policy* 50(1), 102928. DOI: 10.1016/j.telpol.2025.102928. **[2024–2026, EU context]** *Relevance:* Analyses restriction intensity, classification of illegality, and moderation speed in the DSA-TDB to argue that "while the DSA-TDB marks a step towards transparency, significant shortcomings remain: limited usability and accessibility, absence of key data for scrutiny and for monitoring the spread of illegal content, and concerns about consistency, reliability and validity of platform reporting. Consequently, the database falls short of its objectives." Provides the critical normative frame for RQ1's "comparability" sub-question. *Dimensions: 1, 5.* REPEC

Papaevangelou, C., & Votta, F. (2025). "Trading nuance for scale? Platform observability and content governance under the DSA." *Internet Policy Review* 14(3). DOI: 10.14763/2025.3.2037. **[2024–2026, EU context]** *Relevance:* Analyses 439M SoRs from eight VLOPs and finds that >99% of moderation decisions are applied EU/EEA-wide and that automation strongly correlates with faster and territorially uniform enforcement. TikTok, YouTube and X stand out with the most distinct practices — directly relevant for the thesis's TikTok-vs-X focus. *Dimensions: 1, 2, 5, 6.* Policy Review + 2

Tessa, B., Shahi, G. K., Trujillo, A., & Cresci, S. (2026). "When Transparency Falls Short: Auditing Platform Moderation During a High-Stakes Election." arXiv:2604.19285. **[2024–2026, EU context]** *Relevance:* Follow-up that examines whether DSA-TDB inconsistencies persist a year on; specifically documents TikTok-side data quality problems (notably 12,010 SoRs reporting content moderated *before* publication) — a direct methodological warning for the thesis's TikTok longitudinal panel. *Dimensions: 1, 4, 5.* ResearchGate

"Disarranged Harmonization of Transparency Reporting by Social Media Platforms Under the Digital Services Act" (2026). arXiv:2605.17655. *Relevance:* Late-2026 audit of harmonisation between the TDB and the Article 15 transparency reports following the July 2025 schema change; assesses formatting, timeliness, consistency and completeness across platforms and proposes targeted improvements. Directly relevant to the thesis's API v1→v2 mapping problem. *Dimensions: 1, 5.*

Dimension 2 — Automated / Algorithmic Content Moderation

Gorwa, R., Binns, R., & Katzenbach, C. (2020). "Algorithmic content moderation: Technical and political challenges in the automation of platform governance." *Big Data & Society* 7(1). DOI: 10.1177/2053951719897945. *Relevance:* Canonical primer on automated moderation systems.

Argues that algorithmic moderation, even when "well optimised," exacerbates opacity, fairness problems, and the political character of speech decisions — providing the critical lens through which the thesis interprets the `automated_detection` and `automated_decision` flags. *Dimensions: 2, 3, 5.* Sage Journals

Wei, J. T.-Z., Zufall, F., & Jia, R. (2023). "Operationalizing content moderation 'accuracy' in the Digital Services Act." arXiv:2305.09601. *Relevance:* Interdisciplinary refinement of "accuracy" as required by the DSA, mapping it to precision/recall. Useful for the thesis when discussing what platforms actually report under Articles 15(1)(e) and 42, and the limits of comparing automation rates between TikTok and X. *Dimensions: 2, 5.* arxiv

Horta Ribeiro, M., Cheng, J., & West, R. (2023). "Automated Content Moderation Increases Adherence to Community Guidelines." *Proceedings of the ACM Web Conference 2023 (WWW '23)*. arXiv:2210.10454. *Relevance:* Fuzzy regression-discontinuity study on 412M Facebook comments showing causal effects of automated removal on subsequent rule-breaking. Provides empirical justification for treating automation rates as a meaningful governance variable rather than a mere efficiency metric. *Dimensions: 2.* arxiv

Zannettou, S. (2021). "'I Won the Election!': An Empirical Analysis of Soft Moderation Interventions on Twitter." *Proceedings of the 15th International AAAI Conference on Web and Social Media (ICWSM '21)*. arXiv:2101.07183. *Relevance:* One of the first systematic empirical studies of soft moderation (warning labels) on Twitter around the 2020 US election. Highlights inconsistencies in label application — a useful comparison point for X's reported moderation patterns under the DSA. *Dimensions: 2, 4, 6.* Semantic Scholar

Buri, I., & van Hoboken, J. (2021). "The Digital Services Act: Towards a More Accountable Online Public Sphere?" Working Paper, Institute for Information Law (IViR). (Referenced via Buri 2024 in user's list.) *Relevance:* Foundational legal-empirical articulation of accountability obligations in the DSA, with attention to automated decision-making and statement-of-reasons obligations. *Dimensions: 2, 5.*

Dimension 3 — Platform Affordances and Content Governance

Bucher, T., & Helmond, A. (2018). "The Affordances of Social Media Platforms." In J. Burgess, A. Marwick, & T. Poell (eds.), *The SAGE Handbook of Social Media* (pp. 233–253). London: SAGE Publications. (Already in user's list.) *Relevance:* Core theoretical anchor for RQ1; introduces the multi-layered, platform-sensitive view of affordances that the thesis will use to argue that video-vs-text modality matters structurally. *Dimensions: 3.* Digital Academic Repository

Bossetta, M. (2018). "The Digital Architectures of Social Media: Comparing Political Campaigning on Facebook, Twitter, Instagram, and Snapchat in the 2016 U.S. Election." *Journalism & Mass Communication Quarterly* 95(2), 471–496. DOI: 10.1177/1077699018763307. *Relevance:* Operationalises affordance theory into a four-dimension typology (network structure, functionality, algorithmic filtering, datafication). Provides the direct analytical framework for comparing TikTok and X as different "digital architectures" and is essential complementary reading to Bucher & Helmond for RQ1. *Dimensions: 3, 6.*

🟢 **Zeng, J., & Kaye, D. B. V. (2022).** "From content moderation to visibility moderation: A case study of platform governance on TikTok." *Policy & Internet* 14(1). DOI: 10.1002/poi3.287. *Relevance:* Introduces "visibility moderation" as a distinct mode of governance specific to TikTok's video/feed architecture. Crucial for arguing in RQ1 that TikTok's modality structurally favours pre-emptive automated visibility moderation that may be under-captured by the SoR mechanism (which centres on removal-type actions). *Dimensions: 3, 2, 6.* Utrecht University

🟢 **Gray, J. E., & Suzor, N. P. (2020).** "Playing with machines: Using machine learning to understand automated copyright enforcement at scale." *Big Data & Society* 7(1). *Relevance:* Demonstrates the technical asymmetries between text/image and video moderation pipelines — supporting the theoretical claim that modality affects automation reliance, which sits at the heart of RQ1. *Dimensions: 3, 2.* Wiley Online Library

---

Dimension 4 — Event-Driven Moderation Dynamics

🟢 **Pierri, F., Luceri, L., Chen, E., & Ferrara, E. (2023).** "How does Twitter account moderation work? Dynamics of account creation and suspension on Twitter during major geopolitical events." *EPJ Data Science* 12, Article 43. DOI: 10.1140/epjds/s13688-023-00420-7. *Relevance:* Closest analog for RQ2's design. Examines "a large-scale dataset of 270M tweets shared by 16M users in multiple languages" across the 2022 invasion of Ukraine and the 2022 French Presidential election (the Ukraine sub-dataset alone comprises 230,166,962 tweets from 14,995,636 unique users), identifying spikes in account creation/suspension. The thesis can borrow the temporal-anomaly detection methodology. *Dimensions: 4, 2, 6.* arXiv + 3

🟢 **Shahi, G. K., Tessa, B., Trujillo, A., & Cresci, S. (2025).** [See Dimension 1.] *Relevance for D4:* Crucial counter-finding for RQ2: at platform-aggregate scale, the 2024 EP election did *not* produce systematic moderation shifts. The thesis must position its German-snap-election findings against this. *Dimensions: 4, 1, 5.* Icwsm

🟢 **Solovev, K., Drolsbach, C., Demirel, E., & Pröllochs, N. (2026).** "TikTok Rewards Divisive Political Messaging During the 2025 German Federal Election." *EPJ Data Science* (2026). DOI: 10.1140/epjds/s13688-026-00632-7. Preprint: arXiv:2509.10336. **[2024–2026, German context]** *Relevance:* Computational content analysis of N=25,292 TikTok videos posted by German politicians in the run-up to the snap election. Although it studies *content* rather than *moderation*, it provides direct contextual evidence about the German TikTok information environment during the same window the thesis analyses — useful for triangulating moderation patterns against the content that drove them. *Dimensions: 4, 7.* arXiv

🔵 **Scott, M., & Marsh, O. / DFRLab & AlgorithmWatch (2025).** "The Musk Effect: Assessing X's Impact on Germany's Election Discourse." DFRLab (Atlantic Council) & AlgorithmWatch, February 20, 2025. URL: dfrlab.org/2025/02/20/the-musk-effect-xs-impact-on-germanys-election/. **[2024–2026, German context]** *Relevance:* Empirical analysis of X engagement data

(via Meltwater) on German politicians' posts over the last 12 months pre-election. The baseline metric is striking: "the collective views for all posts from AfD politicians was just over 17,000" in the week before Musk's endorsement of Weidel. Finds the AfD-engagement uplift on X is concentrated in a "Musk bump" on Weidel-related posts (largely from English-language accounts) rather than systematic algorithmic favouritism. Directly relevant context for RQ2 on X. *Dimensions: 4, 7, 3.* AlgorithmWatch DFRLab

🔵 **Center für Monitoring, Analyse und Strategie (CeMAS) (2025).** *Autoritäre Strategien im Netz: Analyse und Monitoring digitaler Risiken rund um die Bundestagswahl 2025*. CeMAS, Berlin. URL: cemas.io/publikationen/btw2025/. **[2024–2026, German context]** *Relevance:* Multi-platform monitoring (Telegram, X, TikTok, YouTube, Facebook) of the period 6 Jan – 28 Feb 2025. Documents the AfD's social-media volume — approximately 5,700 videos on TikTok and 44,000 posts on X between November 2024 and February 2025; 6.5% of analysed AfD-related images on X contained AI-generated elements — and at least four Russian influence campaigns (Doppelgänger, Storm-1516, etc.). Vital contextual grounding for RQ2. *Dimensions: 4, 6, 7.*

🔵 **Weinmann, C., Denkovski, O., & Giannaccini, F. / Democracy Reporting International (2025).** *Filtered for You: Algorithmic Bias on TikTok and Instagram in Germany*. DRI Working Paper, April 10, 2025. **[2024–2026, German context]** *Relevance:* Sock-puppet audit (five accounts, 17–21 February 2025) of TikTok and Instagram For You feeds across a dataset of 1,000 videos. Finds "TikTok was more likely than Instagram to recommend political content, and… its algorithm more effectively targeted users based on inferred political orientation"; for the AfD-aligned account, 83.6% of political content was view-aligned on TikTok vs. 81.8% on Instagram. Useful when interpreting differential moderation intensity. *Dimensions: 4, 3, 6, 7.*

🔵 **Katzy-Reinshagen, A., Degeling, M., Barth, S., & Dorn, M. / ISD Germany (2025).** *Wahlkampf im Feed? Wie TikTok mit parteipolitischen Inhalten im Vorfeld der Bundestagswahl 2025 umgeht*. Institute for Strategic Dialogue Germany, February 22, 2025. **[2024–2026, German context]** *Relevance:* TikTok For-You-Page audit conducted pre-election; finds that while most surfaced content was entertainment, "political content from fan pages of the far-right Alternative für Deutschland (AfD) was disproportionately represented in the first political videos shown." Argues that "limited transparency and data access prevent researchers from understanding the extent to which recommender systems are biased." *Dimensions: 4, 3, 5, 7.* ISD

🔵 **Global Witness (2025).** "X and TikTok Algorithms Push Pro-AfD Content to Non-Partisan German Users." Global Witness Investigative Report, February 19, 2025. **[2024–2026, German context]** *Relevance:* Cross-platform sock-puppet study using exactly 3 new accounts per platform (9 total across TikTok, X and Instagram), seeded with equal interest in CDU, AfD, SPD and Greens. Findings: "On TikTok, 78% (28/36) of this recommended party political content was supportive of AfD. On X, 64% (14/22)." Useful triangulation point — combined with Solovev et al. (2026) and DRI, it suggests that *if* moderation did not respond to election dynamics, the algorithmic-recommendation side did. *Dimensions: 4, 6, 7.* Global Witness

---

Dimension 5 — DSA Transparency and Accountability

Urman, A., & Makhortykh, M. (2023). "How transparent are transparency reports? Comparative analysis of transparency reporting across online platforms." *Telecommunications Policy* 47(3), 102477. DOI: 10.1016/j.telpol.2022.102477. *Relevance:* Comparative analysis of pre-DSA Article-15-style reports against the Santa Clara Principles 2.0 across 13 major platforms (incl. TikTok and Twitter). Establishes the cross-platform comparability problem the DSA-TDB was supposed to solve. *Dimensions: 5, 6.* ACM Digital Library

Papaevangelou, C., & Votta, F. (2024). "Content Moderation and Platform Observability in the Digital Services Act." *Tech Policy Press*, March 2024. URL: techpolicy.press/content-moderation-and-platform-observability-in-the-digital-services-act/. **[2024–2026, EU context]** *Relevance:* Policy-focused precursor to the 2025 *Internet Policy Review* paper. Establishes "observability" (vs. transparency) as the analytical frame for evaluating the DSA-TDB. *Dimensions: 5, 1.* Tech Policy Press

European Commission (2024–2025). "How the Digital Services Act Enhances Transparency Online" / DSA Transparency Database documentation. *Shaping Europe's Digital Future* portal, digital-strategy.ec.europa.eu/en/policies/dsa-brings-transparency. **[2024–2026, EU context]** *Relevance:* The Commission's own documentation states that "from 1 July 2025, the content moderation categories and keywords for submitting statements of reasons to the DSA Transparency Database were updated to be aligned with the data categories in the transparency reports, thus making the two transparency data sources comparable." Primary source for the thesis's API v1→v2 mapping. *Dimensions: 5, 1.* European Commission

Bassan, S. (2025). "Mandatory Transparency Reporting Regulation for Content Moderation: A Comparative Analysis Report." SSRN Working Paper, abstract_id 5227491. **[2024–2026]** *Relevance:* Cross-jurisdictional comparison of transparency-reporting regimes; recommends "standardized metrics to improve cross-platform comparisons and enforcement" — direct policy input for RQ1's comparability question. *Dimensions: 5, 6.* SSRN

Heldt, A. et al. / HIIG (2024). "The DSA's Transparency Reports." *HIIG Digital Society Blog*. URL: hiig.de/en/analysis-of-the-dsas-transparency-reports/. **[2024–2026, EU context]** *Relevance:* Qualitative analysis of the first two reporting rounds (2024) of seven VLOPs' DSA transparency reports — Instagram, Facebook, LinkedIn, Pinterest, Snapchat, TikTok and X. Documents heterogeneity, disconnected data points, and the limitations of the pre-template (pre-July 2025) reporting regime. *Dimensions: 5, 6.*

Husovec, M. (2024). *Principles of the Digital Services Act*. Oxford University Press. Plus ongoing "DSA Newsletter" series at husovec.eu/dsa/. — Already in user's list. *Dimensions: 5.* Platform

---

Dimension 6 — Cross-Platform Comparative Moderation

Trujillo, Fagni & Cresci (2025); Drolsbach & Pröllochs (2024); Papaevangelou & Votta (2025); Kaushal et al. (2024). [See Dimension 1.] — Primary cross-platform comparative empirical work using DSA-TDB.

- **Bossetta, M. (2018).** [See Dimension 3.] — Theoretical framework for cross-platform comparison.
- **Cai, J., Patel, A., Naderi, A., & Wohn, D. Y. (2024).** "Content Moderation Justice and Fairness on Social Media: Comparisons Across Different Contexts and Platforms." *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems (CHI EA '24)*. arXiv:2403.06034. *Relevance:* Experimental comparison (200 US users) of perceived justice/fairness of moderation across Reddit and Twitter under different intervention types. Useful for interpreting differential moderation intensities as a normative phenomenon. *Dimensions: 6, 2.*

- **Anonymous Authors (2025).** "Protecting Young Users on Social Media: Evaluating the Effectiveness of Content Moderation and Legal Safeguards on Video Sharing Platforms." arXiv:2505.11160. *Relevance:* Cross-platform comparison of TikTok, YouTube and Instagram moderation via 3,000 video evaluations. Crucially establishes a Unified Harmful Content Framework — methodologically analogous to the thesis's harmonised taxonomy mapping across API v1/v2. *Dimensions: 6, 2, 3.* arxiv

- **Santos, A., Cazzamatta, R., & Napolitano, C. J. (2025).** "Holding platforms accountable in the fight against misinformation: A cross-national analysis of state-established content moderation regulations." *SAGE Journals*. DOI: 10.1177/17480485251348550. *Relevance:* Qualitative comparative analysis of five national misinformation regimes including Germany's NetzDG. Provides the comparative-regulatory backdrop for understanding what DSA territorial-scope filtering means for Germany. *Dimensions: 6, 5.*

Dimension 7 — Agenda-Setting on Social Media

- **Feezell, J. T. (2018).** "Agenda Setting through Social Media: The Importance of Incidental News Exposure and Social Filtering in the Digital Era." *Political Research Quarterly* 71(2), 482–494. DOI: 10.1177/1065912917744895. (Already in user's list.) *Dimensions: 7.*

- **McCombs, M. E., & Shaw, D. L. (1972).** "The Agenda-Setting Function of Mass Media." *Public Opinion Quarterly* 36(2), 176–187. (Founding citation; already in theoretical framework.) *Dimensions: 7.*

- **Gilardi, F., Gessler, T., Kubli, M., & Müller, S. (2022).** "Social Media and Political Agenda Setting." *Political Communication* 39(1), 39–60. DOI: 10.1080/10584609.2021.1910390. *Relevance:* Empirical analysis of Swiss media and Twitter agenda-setting dynamics 2018–2019. Demonstrates that media influence still operates online but in altered, contingent ways — providing a theoretical bridge from Feezell to platform-specific (TikTok/X) agenda-setting in elections. *Dimensions: 7.*

- **Lindskow, K. et al. (2024).** "Algorithmic Agenda-Setting: The Subtle Effects of News Recommender Systems on Political Agendas in the Danish 2022 General Election." *Information, Communication & Society*. DOI: 10.1080/1369118X.2024.2334411. **[2024–2026, EU context]**

Relevance: Empirically extends agenda-setting to algorithmic news curation in a national election context — methodologically parallel to the thesis's snap-election approach (Danish 2022 ≈ German 2025 as algorithmic-agenda-setting cases). *Dimensions: 7, 3.*

🟢 **Kim, R. M., & Anderson, A. (2025).** "The Agenda-Setting Function of Social Media." University of Toronto Working Paper. URL: cs.toronto.edu/~ashton/pubs/agendasetting-2025.pdf. **[2024–2026]** *Relevance:* Recent re-articulation of social-media agenda-setting using HDBSCAN topic-modelling on news production and consumption; provides operational baselines for measuring agenda salience. *Dimensions: 7.*

🟢 ***Scientific Reports* (2025).** "Ideology and polarization set the agenda on social media." *Scientific Reports*. DOI: 10.1038/s41598-025-19776-z. **[2024–2026]** *Relevance:* Empirical evidence that ideological homogeneity (rather than mainstream media) increasingly drives agenda-setting on social media — relevant for interpreting whether moderation intensity correlates with ideologically salient content. *Dimensions: 7.*

PART B — STUDIES GROUPED BY RQ / THEORETICAL FRAMEWORK

RQ1 — Modality, Affordances and Comparability (Bucher & Helmond, 2018)

Core theoretical layer:

- Bucher & Helmond (2018) — *anchor* [D3].
- Bossetta (2018) — operationalises affordance comparison along network/functionality/algorithm/datafication axes [D3, D6].
- Zeng & Kaye (2022) — TikTok-specific visibility moderation [D3, D2, D6].
- Gray & Suzor (2020) — technical asymmetries of video vs. text moderation [D3, D2].

Empirical layer — cross-platform comparison of moderation distributions and automation:

- Drolsbach & Pröllochs (2024) — 350× ratio TikTok-vs-X [D1, D2, D6].
- Trujillo, Fagni & Cresci (2025) — full 100-day platform-wise audit [D1, D2, D5, D6].
- Papaevangelou & Votta (2025) — automation–territorial–speed correlation across VLOPs [D1, D2, D5, D6].
- Kaushal et al. (2024) — ToS-vs-illegality grounds skew (>99.8%) [D1, D5].

Methodological / quality layer — comparability and reporting granularity:

- Tessa, Amram, Monreale & Cresci (2025) — verification processes [D1, D5].
- Groesch et al. (2026) — "Big data small answers" [D1, D5].
- "Disarranged Harmonization" (2026) [D1, D5].
- Urman & Makhortykh (2023) — pre-DSA cross-platform comparability problem [D5, D6].

- HIIG / Heldt et al. (2024) — qualitative review of VLOP reports [D5, D6].

Critical / governance layer:

- Gorwa, Binns & Katzenbach (2020) — political character of automated moderation [D2, D3, D5].
- Husovec (2024) — *Principles of the DSA* [D5].

RQ2 — Event-Driven Moderation and Automation Shifts (Agenda-Setting; McCombs & Shaw, 1972; Feezell, 2018)

Core theoretical layer:

- McCombs & Shaw (1972) — founding article [D7].
- Feezell (2018) — *anchor* for social-media agenda-setting [D7].
- Gilardi et al. (2022) — empirical update on contingent media→public agenda effects [D7].
- Lindskow et al. (2024) — algorithmic agenda-setting in a Danish national election [D7, D3].
- Kim & Anderson (2025); *Sci. Rep.* (2025) — recent agenda-setting under algorithmic curation [D7].

Empirical layer — event-driven moderation dynamics:

- Pierri, Luceri, Chen & Ferrara (2023) — Twitter moderation during geopolitical events (270M tweets) [D4, D2, D6].
- Shahi, Tessa, Trujillo & Cresci (2025) — 1.58B SoRs across eight months around 2024 EP election; negative result on systematic adaptation [D1, D4, D5].
- Tessa, Shahi, Trujillo & Cresci (2026) — follow-up high-stakes election audit [D1, D4, D5].
- Zannettou (2021) — soft moderation around US 2020 election [D2, D4, D6].

German 2025 snap-election contextual layer (grey lit):

- Solovev, Drolsbach, Demirel & Pröllochs (2026 / arXiv 2025) — TikTok content during the German federal election [D4, D7].
- Scott & Marsh / DFRLab & AlgorithmWatch (2025) — "Musk Effect" on X [D4, D7, D3].
- CeMAS (2025) — Bundestagswahl monitoring [D4, D6, D7].
- DRI / Weinmann, Denkovski & Giannaccini (2025) — "Filtered for You" sock-puppet audit [D4, D3, D6, D7].
- ISD Germany / Katzy-Reinshagen et al. (2025) — TikTok FYP audit [D4, D3, D5, D7].
- Global Witness (2025) — cross-platform algorithmic-recommendation study (3 accounts per platform; 78% TikTok / 64% X pro-AfD) [D4, D6, D7].

Critical / regulatory layer:

- European Commission (2024–2025) DSA-TDB & template documentation — primary source on July 2025 schema change [D5, D1].
- Bassan (2025) — comparative transparency-reporting regulation [D5, D6].

Recommendations

1. **Anchor the empirical contribution against Shahi et al. (2025)'s negative finding.** Their conclusion that platforms did *not* significantly adjust moderation around the 2024 EP election (across 1.58B SoRs in eight months) is the natural baseline; the thesis's value lies in testing whether the German snap election — a more intense, national, German-language event with extraordinary far-right activity — produces a different result. Frame RQ2 explicitly as a confirmation/refutation test rather than a free-standing description.
2. **Use Bossetta (2018) plus Zeng & Kaye (2022) to give Bucher & Helmond analytical traction.** The Bucher & Helmond chapter is theoretical; combining it with Bossetta's typology and Zeng & Kaye's "visibility moderation" concept will let RQ1 do real comparative work between TikTok (video, visibility-moderation-heavy, high SoR volume) and X (text, low-volume reporting, post-Musk policy regime).
3. **Treat the July 2025 API v1→v2 break as a substantive finding, not just a methodological hurdle.** No peer-reviewed paper yet evaluates it empirically — citing the EC documentation (which states the alignment took effect 1 July 2025), Husovec's newsletters, and the "Disarranged Harmonization" arXiv preprint (2026), the thesis can position the schema break itself as evidence of regulatory learning, and probe whether reporting granularity dropped or rose after July 2025.
4. **Triangulate moderation findings with the German grey literature.** CeMAS (5,700 AfD TikTok videos, 44,000 AfD X posts; 6.5% AI-generated AfD images on X), DRI (83.6% view-aligned political content on TikTok for AfD-leaning users), ISD Germany, DFRLab/AlgorithmWatch ("just over 17,000" total weekly AfD views on X before the Musk bump), and Global Witness (78% TikTok / 64% X pro-AfD recommendations) all directly observed TikTok and X during the same window. If the DSA-TDB shows little moderation response (per the Shahi et al. analog), grey-literature evidence of pro-AfD algorithmic amplification becomes the agenda-setting "second story" the thesis should foreground.
5. **Cite Kaushal et al. (2024) and Groesch et al. (2026) whenever interpreting decision-grounds distributions.** The fact that >99.8% of SoRs are ToS-based (not illegality-based) means the "removal category distribution" in RQ1 is largely a window onto platforms' own terms — not into illegal content per se. Be explicit about this in the methods section.
6. **Benchmark thresholds for changing recommendations:** if the thesis observes (a) automation rates jumping above ~95% on either platform during the election window, or (b) a measurable drop in optional-field completeness (`content_language`, `decision_visibility_other`, free-text grounds) after the API v2 switchover, treat these as findings that warrant stronger regulatory recommendations in the discussion section.

Caveats

- **arXiv-status sources** (Tessa et al. 2026; Shahi et al. 2025 ICWSM workshop version; "Disarranged Harmonization" 2026) are preprints/workshop versions — final peer-reviewed venues may differ; cite the arXiv ID and check for updates before submission. The Solovev et al. paper has now been formally published in *EPJ Data Science* (DOI: 10.1140/epjds/s13688-026-00632-7), superseding the arXiv:2509.10336 preprint.
- **Civil-society sock-puppet studies** (DRI: 5 accounts / 1,000 videos / 5-day window; Global Witness: exactly 3 accounts per platform / 9 total; ISD Germany) use small Ns over short windows; cite as grey-literature corroboration, not as primary empirical evidence.
- **The thesis explicitly excludes the 2024 EP election as an analog.** Shahi et al. (2025) is therefore included as a methodological/null-result benchmark, not as a substantive analog — be careful to maintain that framing.
- **Drolsbach & Pröllochs (2024) is a short workshop paper** (4 pp.), descriptive rather than inferential. Its TikTok-vs-X ratio (>350×) is widely cited but should be treated as descriptive, not causal.
- **The Shahi et al. (2025) abstract states "eight months"** of observation while the paper body refers to "ten months"; cite the abstract version and note this internal inconsistency if relevant.
- **Some included sources (e.g., "Protecting Young Users on Social Media" arXiv:2505.11160) appear under author-anonymous arXiv format** as of the search date; check arXiv for the final author list before citation.